

AI Automation & WhatsApp Bot

Company Knowledge Base • RAG / MCP / LLM Optimized

Purpose: authoritative company knowledge for RAG retrieval, MCP-connected agents, and LLM responses. Sections are intentionally self-contained so retrieval can return precise context without requiring the model to infer missing business facts.

01. Service Overview

What we provide

We design AI-powered automation systems that connect conversational AI with business workflows, knowledge bases, APIs, databases, and messaging channels such as WhatsApp.

Primary outcomes include faster customer responses, automated FAQs, lead qualification, customer support automation, internal knowledge access, structured data capture, and reduced repetitive manual work.

Typical use cases

FAQ assistant, product/service information, lead capture, lead qualification, customer support, appointment or inquiry workflows, order/status lookup when connected to a backend, human handoff, conversation summaries, and workflow automation.

02. WhatsApp Bot Architecture

Typical flow

User message → WhatsApp/channel layer → webhook/API → orchestration layer → context and intent handling → RAG retrieval when required → LLM generation → authorized business tools/APIs → response → logging and analytics.

Tool access

MCP can provide a controlled interface between the AI agent and approved tools or services. Tools should expose only the actions required by the business and should enforce authentication, authorization, validation, and safe error handling.

03. RAG Knowledge Base

Knowledge design

Company knowledge should use clear headings, stable terminology, explicit service names, FAQs, policies, constraints, and structured facts.

Chunk information by semantic meaning rather than arbitrary page length. Preserve metadata such as topic, service, audience, source, version, and language where useful.

Use retrieved company context for company-specific answers. The model should not invent prices, guarantees, policies, capabilities, availability, customer records, or commitments.

Retrieval quality

Good retrieval depends on meaningful chunking, accurate embeddings or retrieval indexes, metadata, query rewriting when needed, consistent terminology, and high-quality source content.

An English source document can still support multilingual responses. Source-language duplication is not mandatory for every output language.

04. Multilingual Response Policy

Supported conversation styles

The bot may answer in English, Hindi written in Latin script (Hinglish), or Hindi written in Devanagari.

Default behavior: answer in the language and script used by the user's latest message unless the user explicitly requests another language.

Keep company names, product names, SKUs, URLs, API identifiers, code terms, and exact technical field names unchanged when translating them would reduce accuracy.

Why English knowledge can work

RAG retrieves semantic meaning rather than simply matching identical words. A capable multilingual embedding/retrieval layer and multilingual LLM can connect a Hindi or Hinglish question to relevant English knowledge.

For maximum reliability, important company terminology can include aliases and multilingual keywords in metadata or FAQ entries. This is an optimization, not a requirement to duplicate the whole knowledge base in three languages.

05. Groq / LLM Integration

Provider role

Groq can be used as the LLM inference provider. The LLM generates the final response; the RAG layer supplies company-specific context; MCP or application tools provide controlled actions.

The provider can be changed later without rewriting the company knowledge base if the application keeps retrieval, business logic, and provider integration separated.

API key security

Keep the Groq API key server-side. Store it in environment variables or a secure secrets manager. Never expose it in browser JavaScript, public source code, client-side configuration, or a frontend bundle.

06. Bot Safety & Fallback

Unknown information

If retrieved knowledge does not contain the requested fact, the bot should clearly state that the information is not available in its current company knowledge and, where appropriate, offer human assistance or ask for clarification.

High-risk actions

Do not fabricate order status, prices, discounts, delivery commitments, legal claims, credentials, or internal policies. Irreversible or sensitive business actions should require appropriate authorization and confirmation.

07. RAG Retrieval Facts

Structured facts

service=ai-automation

channel=whatsapp

architecture=webhook,orchestration,RAG,LLM,MCP,tools,logging

capabilities=FAQ,lead-capture,qualification,support,knowledge-search,human-handoff,workflow-automation,multilingual

ai-stack=RAG,MCP,LLM,Groq

languages=English,Hinglish-Latin,Hindi-Devanagari

security=server-side-api-key,environment-secrets,authorization,validation

fallback=no-invention,human-handoff,clarification